

Chapter 14

Agent-based Model in MATLAB

Agent-based Computational Economics is one of the newer fields in economics. Agent-based models simulate the behavior of multiple heterogeneous agents interacting in a variety of ways. While the modeling of economic agents has a long tradition in economics, agent-based modeling departs from it in a number of ways. For example, when modeling a market economy, the standard neoclassical competitive general equilibrium approach usually assumes that agents have fixed preferences, perfect and complete information, no reproductive behavior, and also that trade is organized by a central auctioneer that given all agents preferences and endowments computes the set of equilibrium prices. Thus, agents are price-takers and do not engage in trade at prices other than those given by the central auctioneer. Also space, that is geography, is usually an absent dimension in that approach. In contrast, agent-based models allow agents to display a number of more realistic characteristics and behaviors, i.e. changing preferences, bounded rationality and memory, imperfect and incomplete information, and local trade - agents may interact with neighbors in a geographically defined space and prices emerge from these decentralized interactions.

In this chapter we will introduce a famous agent-based model known as the Sugarscape model, developed by Joshua M. Epstein and Robert Axtell (1996). This is a model designed to simulate a variety of social phenomena such as population dynamics, migration, interaction with the environment, trade, group formation, combat and transmission of culture. We will learn how to represent and simulate the simplest version of this model in MATLAB. To do this, the knowledge of basic MATLAB operations and data types - vectors and matrices, with the addition of data types named structures and cell arrays that we will explain below - will suffice. However, more sophisticated simulations may require the use of object oriented programming techniques, something also available in MATLAB - see “MATLAB Classes and Objects” in the “Programming and Data Types” section of the MATLAB help navigator - as well as in lower level object oriented programming languages such as C++, C# or Java.

1. The Sugarscape Model: Introduction

The version of the classic Sugarscape model that we use in this chapter can be thought of as two major cities located near one another like Dallas and Fort Worth in Texas or Minneapolis and St. Paul in Minnesota. There is an original distribution of stores of a certain type in this terrain; for example, coffee houses such as Starbucks or perhaps mailing and business services stores such as UPS Stores. The franchise owners at each location work with varying degrees of efficiency and thus have different costs. They thus require different levels of revenues in order to continue to make a profit. Their profit each period is added to their accumulated wealth; however, if this wealth goes to zero the franchise is shut down. The surviving franchise owners each period look around for a nearby location that would be more favorable and move the store if they find a higher revenue location. However, some of the franchise owners scout longer distances away from their present store than others.

More formally, the Sugarscape model consists of two main elements: a terrain where events unfold named “sugarscape”, which contains the spatial distribution of a generalized resource named “sugar” which can be thought of as the customer potential or revenue level at that location. The agents have metabolism levels and must eat to survive. This metabolism may be thought of as the cost of running the business in each period. Thus the difference between the sugar that the agents obtain at their location in each period and their metabolism level is like the profit of the enterprise in each period. This profit is accumulated as wealth from period to period; however, if the wealth level goes to zero in any period the agent dies, i.e. goes out of business. Thus, the agents are characterized by a set of fixed states (genetic characteristics such as metabolism and length of vision) and variable states (such as location and wealth) and move around the sugarscape following simple rules of behavior.

The sugarscape is represented by a two-dimensional coordinate grid or lattice. At every point of the grid given by the coordinates (x,y) there is a sugar level. Thus, we can easily represent the sugarscape in MATLAB by means of a matrix. For example, if we want to create and display a (50×50) sugarscape with a level of sugar equal to 4 units in the southwest quadrant and a level of 2 units elsewhere, we can do it with the following statements

```

for i = 1: 50;
    for j = 1: 50;
        if (i > 25 & j < 25)
            s(i,j) = 4;
        else
            s(i,j) = 2;
        end
    end
end
image(s);

```

In the statements above `image(s)` is a MATLAB function that displays the array `s`.

Figure 14.1 below shows the result, where the lighter region corresponds to the value 4 and the darker region corresponds to the value 2.




Figure 14.1 Sugarscape with Two Levels of Sugar

To represent agents, we can use another data type available in MATLAB called a structure. A structure is an array with “data containers” named “fields”. These fields can contain any kind of data. For example, let’s assume that every agent is characterized by two states: active, which signals if the agent is alive or not, with values equal to 1 and 0 respectively, and metabolism, that is the amount of sugar each agent has to eat per time period to survive. The statements

```
a_str.active = 1;
a_str.metabolism = 4;
```

create the simple 1x1 structure `a_str` containing two fields. If we use the statements

```
a_str(2).active = 1;
a_str(2).metabolism = 3;
```

then `a_str` becomes a 1x2 array with two fields. Let's assume that we want to create and display a random population of agents - say all those for whom the corresponding value from a [0,1] uniform distribution is lower than 0.2 - on a 50x50 grid. Also, we will assume that there can only be one agent on each location. We can achieve this with the following statements

```
for i = 1:50;
    for j = 1:50;
        if (rand < 0.2)
            a_str(i,j).active = 1; %put an agent on this location
            a_str(i,j).metabolism = 3;
        else
            a_str(i,j).active = 0; %keep this location empty
            a_str(i,j).metabolism = 0;
        end
    end
end
```

With these statements we can create a structure with 2,500 elements, each with two fields.

If we want to display the location of every agent on the grid, we can do it with the following statements, where we transfer the elements of the field `active` into the `a` matrix, and where the MATLAB function `spy(a)` displays all the nonzero elements in matrix `a`.

```
for i = 1:50;
    for j = 1:50;
        a(i,j) = a_str(i,j).active;
    end
end
spy(a);
```

The result, with a number of agents equal to 474, is shown in Figure 14.2 below, where `nz` means the number of non-zero elements.




Figure 14.2 Agents Locations

Now that we have introduced the basic building blocks of the Sugarscape model and its MATLAB representation, we can move on to a more detailed presentation.

2. The Sugarscape Model

Next we present a more complex topography for the sugarscape and also more complex agent characteristics. We will also define rules that will govern the autonomous growth of sugar in the sugarscape and the movement of the agents on it.

We will assume that the sugarscape is characterized by two mountains of sugar, one in the southeast portion of the grid, and the other in the northwest, and that these two mountains are symmetric. Thus, for a 50x50 grid, we will assume that one peak of the sugarscape is approximately on the $(0.75 * 50, 0.25 * 50)$ coordinate, while the other is on the $(0.25 * 50, 0.75 * 50)$ coordinate. From the peaks down, the level of sugar at each location will follow decreasing paths.

We will also specify a very simple growback rule for the sugarscape:

Sugarscape rule G_{∞} : Grow back to full capacity immediately.

Thus, at each run of the model, the level of sugar grows back to its initial level. The symbol G_{∞} here is a fancy way to specify how rapidly the amount of sugar (revenue) grows back in each time period. Epstein and Axtell (1996) use a variety of such rules.

We will also assume that the sugarscape is what in geometry is known as a Torus, or in a more familiar way, that it corresponds to the surface of a donut. This means, for example, that an agent moving to the south on column 6, after reaching row 50 will appear on the sugarscape from the north in the coordinate (1,6), and an agent moving to the east on row 6, after reaching column 50 will appear on the sugarscape from the west on the coordinate (6,1). Analogous patterns will be followed by agents moving north or west.

Turning now to the agents, we will assume that each agent has four characteristics, two of them fixed and the other two variable. The fixed ones are metabolism - the amount of sugar the agent has to consume at each time period to stay alive - and vision - the number of sites in the sugarscape each agent can see. We will assume that agents can see only in four directions: north, south, east and west. Thus, they can not see in diagonal directions. The level of vision is the maximum number of sites each agent can see in a given direction. Metabolism and vision are genetic characteristics randomly distributed among agents.

The variable characteristics of agents are location on the sugarscape and wealth, with the latter understood as the agent's stock of sugar. We will assume that agents are randomly born around the sugarscape at the beginning of the simulation. Each agent will start its life with a level of wealth equal to the level of sugar in the sugarscape location where it was born.

We will specify a rule that will govern the behavior of each agent on the sugarscape:

Agent movement rule M:

- Look out as far as vision permits in the four principal directions and identify the unoccupied site(s) having the most sugar
- If the greatest sugar value appears on multiple sites then select the nearest one
- Move to this site
- Collect all the sugar at this new position

Once sugar is collected, the agent's wealth is incremented by the sugar collected and decremented by its metabolic rate. An agent lives forever, unless its wealth is below its metabolic rate. In this case, it dies and is removed from the sugarscape. In principle,

all agents should apply this rule simultaneously. However, since the simulation is run on a serial computer, only one agent will be active at any instant. In this case, it is recommended to randomize agents' order of movement, and we will do this in the MATLAB code. We will also randomize step one of the rule, that is, the order in which each agent searches the four directions.

Having presented the building blocks of the simplest version of the Sugarscape model, we now turn to its MATLAB representation.

3. The Sugarscape Model in MATLAB

The MATLAB representation consists of a main program named `sugarscape1.m` and a number of functions, all of which are available from the book web site. Below is the code of the main program.

```

%Initialize model parameters
nruns = 6;
size = 50; %even number
metabolismv = 4;
visionv = 6; %set always smaller than size
maxsugar = 20;

%Initialize sugarscape and display
s = initsugarscape(nruns, size, maxsugar);

%Initialize agents population
a_str = initagents(size, s, visionv, metabolismv);

%Main loop (runs)
for runs = 1:nruns;
    % Display agents' locations
    dispagentloc(a_str, size, nruns, runs);

    % Select agents in a random order and move around the sugarscape %
    % following rule M
    for i = randperm(size);
        for j = randperm(size);
            if (a_str(i,j).active == 1) %is there an agent on this
                                    %location?
                %Agent explores sugarscape in random directions and
                %selects best location
                temps = s(i,j);
                tempi = i;
                tempj = j;
                for k = a_str(i,j).vision : -1 : 1;
                    [temps, tempi, tempj] =
                        see(i,j,k,a_str,s,size,temps,tempi,tempj);
                end
                %Agent moves to best location, updates sugar stock and
                %eats sugar
                a_str = moveagent(a_str, s, i, j, temps, tempi, tempj);
            end % if
        end % for j
    end % for i
end % for runs

```

The program begins with the initialization of the model parameters - the number of runs, the size of the sugarscape, the maximum value of metabolism and vision of the agents, and the maximum level of sugar in the sugarscape. Then follows a call to the

function named `initsugarscape`, which will return a matrix named `s` containing the sugar levels in the sugarscape. Next a call to the function `initagents` returns the data structure `a_str` which will contain the agents' population.

Then follows the main loop of the program corresponding to the number of runs - each run represents a time period - of the simulation. At each pass of the loop, the locations of the agents on the sugarscape are displayed as a way of visualizing their movements. This is achieved by calling the function `dispagentloc`.

Then each agent, in a random order, explores the sugarscape, selects the best location, updates its wealth and eats sugar to survive. This section of the program begins with the following statements.

```
for i = randperm(size);
    for j = randperm(size);
```

The `randperm(n)` function performs a random permutation of the elements of the set $(1, 2, \dots, n)$. Thus the `randperm(size)` MATLAB function creates a vector with a number of elements equal to `size` and performs a random permutation of those elements. Thus, once the two `for` loops - one for `i` and the other for `j` - are completed, the whole population of agents will have moved but in a random order. The conditional

```
if (a_str(i,j).active == 1) %is there an agent on this location?
```

checks if there is an active agent in the (i, j) location being examined, where a 1 in the field `active` of the agent data structure denotes that there is an agent, while a 0 denotes the opposite. Then, if there is an active agent in the location, the program proceeds to apply the agent's rule of movement, while if that is not the case it proceed to examine another location looking for an active agent. The agent's rule of movement is implemented with the statements below

```

%Agent explores sugarscape in random directions and
%selects best location
temps = s(i,j);
tempi = i;
tempj = j;

for k = a_str(i,j).vision : -1 : 1;
    [temps, tempi, tempj] =
        see(i,j,k,a_str,s,size,temps,tempi,tempj);
end

%Agent moves to best location, updates sugar stock and
%eats sugar
a_str = moveagent(a_str, s, i, j, temps, tempi, tempj);

```

The statements begin with the setting of three temporary variables. The variable `temps` contains the level of sugar in the agent's current location, while `tempi` and `tempj` contain the location's coordinates. Then follows a loop that goes from the agent's maximum level of vision to 1, in decrements of one unit. At each pass of this loop, the function `see` is called. This function will see around the agent's neighborhood in the north, south, east and west directions, from the farthest position the agent can see to its immediate surroundings, and will return the maximum level of sugar in the variable `temps` and its location coordinates in the variables `tempi` and `tempj` respectively. Finally, once the loop is completed, the function `moveagent` is called to move the agent to the new location and to update its stock of wealth.

From this overview of the main program we turn next to descriptions of the functions.

3. Functions

3.1 Initsugarscape

The “initsugarscape” function initializes the level of sugar at each location of the sugarscape. To better understand the procedure used, we will begin with simpler examples. Suppose that we want to generate an 11x11 sugarscape `s1` with a single mountain with a peak in the center. The corresponding statements are shown below, where `i` and `j` are the matrix coordinates varying from 1 to 11. The vectors `x` and `y` are two identical eleven-element vectors containing the values `[-5 -4 -3 -2 -1 0 1 2 3 4 5]`.

```

%Generate sugarscape with one peak in the center
x = -5:5;
y = -5:5;
maxsugar = 20;
for i = 1:11;
    for j = 1:11;
        if (x(i) == 0 & y(j) == 0)
            s1(i,j) = maxsugar;
        else
            s1(i,j) = maxsugar / (abs(x(i)) + abs(y(j)));
        end
    end
end
end

```

The value of each element in the `s1` matrix is computed dividing the given maximum level of sugar by the sum of the absolute value of the corresponding elements in the `x` and `y` vectors as shown below:

$$s1(i,j) = \text{maxsugar} / (\text{abs}(x(i)) + \text{abs}(y(j)));$$

where `abs` is the absolute value. The peak of the mountain will be where the corresponding elements of the `x` and `y` vectors equal zero. Thus, the value of `s1(6,6)`, which will be located at the center of the sugarscape, will be equal to `maxsugar` - making a minor adjustment to avoid the division by zero. And the values on the corners - i.e. `s1(1,1)` - will be equal to `(maxsugar/10)`. All the other values would be, in a decreasing order, between `maxsugar` and `(maxsugar/10)` as shown in Figure 14.3 below.




Figure 14.3 Sugarscape with a Center Peak

Now, if we want to generate a sugarscape with a peak in the southeast instead of the center, the values of x and y should be shifted to

$$x = [-9 \ -8 \ -7 \ -6 \ -5 \ -4 \ -3 \ -2 \ -1 \ 0 \ 1]$$

and

$$y = [-3 \ -2 \ -1 \ 0 \ 1 \ 2 \ 3 \ 4 \ 5 \ 6 \ 7].$$

In this case, the peak of the sugarscape will be in the $s1(10, 4)$ location, as shown in Figure 14.4 below.




Figure 14.4 Sugarscape with a South-West Peak

The `initsugarscape` function initializes the level of sugar at each location of the sugarscape. This particular function will generate a topography characterized by two mountains of sugar, one in the southwest portion of the grid, and the other in the northeast. These two mountains are symmetric. From the peaks down, the level of sugar will follow decreasing paths. The function code is available in file `initsugarscape.m`,

This function begins by generating a sugarscape $s1$ containing a single peak in the southwest. To do so, the “Generate sugarscape with one southwest peak” section of the function, reproduced below, applies a similar procedure to the one just described.

```

%Generate sugarscape with one south west peak
x = -ceil(0.75*size) : size-ceil(0.75*size)-1;
y = -ceil(0.25*size) : size-ceil(0.25*size)-1;

for i = 1:size;
    for j = 1:size;
        if (x(i) == 0 & y(j) == 0)
            s1(i,j) = maxsugar;
        else
            s1(i,j) = maxsugar / (abs(x(i)) + abs(y(j)));
        end
    end
end
end

```

For example, for a value of `size` equal to 50, it begins by generating a 50-element vector `x`. The statement

```
x = -ceil(0.75*size) : size - ceil(0.75*size) - 1;
```

is used to create a 50 element vector of integers as follows. The values in the vector begin at minus the ceiling of the product ($0.75 * 50$), i.e. the next integer above 37.5, namely -38. They end at the value $(50 - 38 - 1)$, i.e. 11. So `x` will be a 50 element vector with the values

$$[-38, -37, \dots, -1, 0, 1, \dots, 10, 11]$$

Thus, the value zero will be in the 39th position of the `x` vector. In a similar way the vector `y`, which goes from -13 to 36, is generated with the value zero in its 14th position.

After doing this, each element of the sugarscape matrix `s1` is generated. The result will be a sugarscape with a peak in the `s1(39,14)` location, i.e. in the southwest corner of the array.

Once the first mountain is generated, a symmetric one is obtained by transposing the matrix `s1` with the statement

```
s2 = s1';
```

Then, the statement

```
s = s1 + s2;
```

generates the two-peak sugarscape. The following two statements

```
maxrow = max(s);
```

```
max(maxrow)
```

compute the row containing the maximum value in the matrix *s* and print the maximum value in this row. This may seem redundant, since we set the parameter `maxsugar` at the beginning of the program. That value is indeed the maximum for the peaks in *s1* and *s2*. But the peaks in *s* will be a bit higher since to each original peak we will be adding the value of the corresponding cell in the symmetric matrix, which will be a low value given its distance from the peak.

The final statements below display the image of the sugarscape shown in Fig.

14.5.

```
figure(1);
```

```
imagesc(s);
```

```
axis square;
```

The statement `figure(1)` generates a figure where an image will be displayed. The statement `imagesc(s)` scales the data in matrix *s* to the full range of colors and displays the corresponding image of the sugarscape matrix *s*. Finally, the statement `axis square` makes the image square. The result is the figure with two centers of economic activity as shown below.




Figure 14.5 Two-peak Sugarscape

Next we turn from the code for the sugarscape to the code for the agents.

3.2 Initagents

The function “initagents” generates a random initial population of agents. Its code is shown below.

```
function a_str = initagents(size, s, visionv, metabolismv);
for i = 1:size;
    for j = 1:size;
        if (rand < 0.2)
            a_str(i,j).active = 1; %put an agent on this location
            a_str(i,j).metabolism = ceil(rand * metabolismv);
            a_str(i,j).vision = ceil(rand * visionv);
            a_str(i,j).wealth = s(i,j);
        else
            a_str(i,j).active = 0; %keep this location empty
            a_str(i,j).metabolism = 0;
            a_str(i,j).vision = 0;
            a_str(i,j).wealth = 0;
        end
    end
end
end
```

The information about agents is stored in the data structure `a_str` with four fields. The field `active` contains a 1 or 0 depending of the situation of the agent in a specific location (active, that is alive; or inactive, that is dead). A location with an inactive agent is treated in the main program and other functions as an empty location. If the values generated by the uniform distribution MATLAB function `rand` are below 0.2, an agent is born.

The fields `metabolism` and `vision` contain the corresponding integers randomly distributed between 1 and the maximum level of each characteristic. The MATLAB function `ceil` is used to round the randomly created `vision` and `metabolism` variables up to the next integer. The field `wealth` is initialized as equal to the amount of sugar in the location of the sugarscape where the agent was born.

3.3 Dispagentloc (display agent location)

This simple function transforms the field `agent` from the `agents` data structure into a matrix named `a` and displays agents' locations, since MATLAB does not allow one to display that field directly. The code of the function is shown below.

```
function a = dispagentloc(a_str, size, nruns, runs);
for i = 1:size;
    for j = 1:size;
        a(i,j) = a_str(i,j).active;
    end
end
figure(2);
subplot(ceil(sqrt(nruns)),ceil(sqrt(nruns)),runs), spy(a);
axis square;
```

The statement `figure(2)` tells MATLAB to display a second figure with the agent's locations - remember that a first figure was created before to display the sugarscape.

Consider next the line of code

```
subplot(ceil(sqrt(nruns)),ceil(sqrt(nruns)),runs), spy(s);
```

and notice that this one line contains two separate MATLAB statements, i.e. the function calls

```
subplot()
and
spy()
```

The call to `subplot` divides the window into a number of panes and the call to `spy` plots the active pane. These statements thus allow us to display multiple images in a single figure such as the images of agents' locations in successive runs of the program. The MATLAB function

```
subplot(m,n,p);
```

creates an axes in the `p`th pane of a figure divided into an `m`-by-`n` matrix of rectangular panes. For example, if we set the number of runs parameter in the main program equal to 8, then the statement


```
subplot(ceil(sqrt(nruns)),ceil(sqrt(nruns)),nruns), spy(s);
```

where `ceil(sqrt(nruns))` is the ceiling (i.e. the integer above) the square root of the number of runs, will divide the figure (window) into a matrix with 3 rows and 3 columns of panes to accommodate the images of the agent's locations in successive runs.

3.4 See and Neighbor

The `see` and `neighbor` functions explore the neighborhood an agent can see according to its level of vision in four directions - north, south, east and west - each direction selected in a random order. Remember that the location coordinates of the agent are given by (i, j) and that the agent's level of vision is equal to k . For each integer between k and 1 - that is, going from the outermost part of the neighborhood to its center - the function will check the level of sugar in each of the four directions. Every time the level of sugar in a location being examined is greater than the level of sugar in the agent's location, the level and coordinates of the higher value found will be stored in the temporary variables `temps`, `tempi` and `tempj` respectively. Thus, at the end of the exploration, these variables will contain the highest level of sugar found and its location.

Imagine that we begin by exploring the neighborhood in the south direction for a level of vision equal to k and from the location (i, j) . Thus, we want to examine the location $(i+k, j)$. If $(i+k \leq \text{size})$, where `size` is the dimension of the sugarscape, there is no problem. However if $(i+k > \text{size})$, we have to remember that in Section 2 above we define the sugarscape as a Torus. Then, in this case the location to be examined will be $(i+k-\text{size}, j)$. For example, if we start from the location $(48, 2)$ with $k = 6$, then the location to be examined will be $(4, 2)$. Thus, to summarize, we could write the following pseudo code, where `neighbor` will be a function that will check the level of sugar in the location (u, v) .

```
if (i + k > size)
    u = i + k - size;
    v = j;
    neighbor(u,v);
else
    u = i + k;
    v = j;
    neighbor(u,v);
end
```

Now, in the case when we want to examine the north direction, the code should be

```
if (i - k < 1) %or equivalently if(k - i > -1)
    u = i - k + size;
    v = j;
    neighbor(u,v);
else
    u = i - k;
    v = j;
    neighbor(u,v);
end
```

Analogous codes could be written for the cases of the east and west directions. However, we want to write a general code encompassing all the four cases. That is, something of the form

```
if ( (1) > (2) )
    u = (3);
    v = (4);
    neighbor(u,v);
else
    u = (5);
    v = (6);
    neighbor(u,v);
end
```

To do so, we proceed as follows. We define the following four vectors, each with six elements:

```
south = [i+k    size    i+k-size    j            i+k    j];
north  = [k-i    -1     i-k+size    j            i-k    j];
east   = [j+k    size    i            j+k-size    i    j+k];
west   = [k-j    -1     i            j-k+size    i    j-k];
```

Next we make use of a MATLAB object named “cell array”. A cell array is an array whose elements are also arrays. For our case, think of it as a matrix whose elements are vectors instead of numbers. The following statements create a cell array of dimension 1x4 whose elements are the vectors south, north, east and west. Notice that the indexes of a cell array are between braces.

```
c{1} = south; c{2} = north; c{3} = east; c{4} = west;
```

Now, for example, if we want to access the third element of the north vector, we can do it using a double indexing notation such as

```
c{2}(3);
```

Then, a general code to explore the neighborhood of an agent, selecting four directions of search in a random manner, can be written as:

```
for m = randperm(4);
    if (c{m}(1) > c{m}(2))
        u = c{m}(3);
        v = c{m}(4);
        [temps, tempi, tempj] =
            neighbor(u,v,a_str,s,temps,tempi,tempj);
    else
        u = c{m}(5);
        v = c{m}(6);
        [temps, tempi, tempj] =
            neighbor(u,v,a_str,s,temps,tempi,tempj);
    end
end
```

To check this go through the south and then the north cases and you should get the same results as those shown above.

We turn now to explain the workings of the `neighbor` function, which is a very simple one. As can be seen in the code above this function receives as inputs, among other arguments, the variables `temps`, `tempi`, and `tempj` and returns the same variables as outputs. Remember that `temps` contains the level of sugar in a given location and `tempi` and `tempj` contain the coordinates of the location. The code of the `neighbor` function is shown below.

```

function [temps, tempi, tempj] =
neighbor(u,v,a_str,s,temps,tempi,tempj);

if (a_str(u,v).active == 0)
    if (s(u,v) >= temps)
        temps = s(u,v);
        tempi = u;
        tempj = v;
    end
end

```

Thus, the function first checks whether the (u, v) location is free so that an agent can move there. If that is the case, it checks to see whether or not the level of sugar in the (u, v) location of the sugarscape is greater than or equal to the one previously found and stored in the variable `temps`. If so, it puts the new level found in the `temps` variable, and its corresponding (u, v) coordinates in the variables `tempi` and `tempj`.

To conclude this section, we reproduce below the entire code of the `see` function.

```

function [temps, tempi, tempj] =
see(i,j,k,a_str,s,size,temps,tempi,tempj);

south = [i+k size i+k-size j i+k j];
north = [k-i -1 i-k+size j i-k j];
east = [j+k size i j+k-size i j+k];
west = [k-j -1 i j-k+size i j-k];

c{1} = south; c{2} = north; c{3} = east; c{4} = west;

for m = randperm(4);
    if (c{m}(1) > c{m}(2))
        u = c{m}(3);
        v = c{m}(4);
        [temps, tempi, tempj] =
            neighbor(u,v,a_str,s,temps,tempi,tempj);
    else
        u = c{m}(5);
        v = c{m}(6);
        [temps, tempi, tempj] =
            neighbor(u,v,a_str,s,temps,tempi,tempj);
    end
end

```

3.5 Moveagent

Once the neighborhood of the agent has been examined, it is time to move the agent to the best location found, update its wealth and let it eat sugar. This is what the `moveagent` function shown below does.

```
function a_str = moveagent(a_str, s, i, j, temps, tempi, tempj);

if (temps > s(i,j))
    % Agent moves to best location and updates wealth
    a_str(tempi,tempj) = a_str(i,j);
    %Set old location to unoccupied
    a_str(i,j).active = 0;
    a_str(i,j).vision = 0;
    a_str(i,j).metabolism = 0;
    a_str(i,j).wealth = 0;
    % update wealth at new location
    a_str(tempi,tempj).wealth = a_str(tempi,tempj).wealth + temps -
                                a_str(tempi,tempj).metabolism;
    % if wealth is less than zero set location to unoccupied
    if (a_str(tempi,tempj).wealth <= 0)
        a_str(tempi,tempj).active = 0;
        a_str(tempi,tempj).vision = 0;
        a_str(tempi,tempj).metabolism = 0;
        a_str(tempi,tempj).wealth = 0;
    end
else
    % Agent stays in position and updates wealth
    a_str(i,j).wealth = a_str(i,j).wealth + temps -
                        a_str(i,j).metabolism;
    if (a_str(i,j).wealth <= 0)
        a_str(i,j).active = 0;
        a_str(i,j).vision = 0;
        a_str(i,j).metabolism = 0;
        a_str(i,j).wealth = 0;
    end
end
end
```

If a new and better location than the one previously occupied by the agent is found, that is, if the statement below is true

```
if (temps > s(i,j))
```

then the agent moves to the new location whose coordinates are stored in the variables `temp_i` and `temp_j`. The old location is set to unoccupied, and the agent's wealth is updated adding to its previous wealth the amount of sugar found in the new location and subtracting the sugar to be consumed according to its metabolic rate. If the resulting level of wealth is less or equal than zero then the agent dies and all its fields are set to zero.

In the case that no better location was found, the agent stays into place, updates its wealth and eats sugar. Again, if the resulting level of wealth is less or equal to zero, the agent dies.

4. Results

We are now ready to analyze the behavior of the population of agents in the sugarscape given the topography, the growback rule G_∞ and the agents' rule of movement M . The agents' locations for six successive runs, for a maximum vision of 6 and a maximum metabolism equal to 4, are shown in Figure 14.6 below. The order of graphs corresponding to the successive runs goes from left to right then down to the next row.




Figure 14.6 Agents' Locations for Six Runs

We can observe that in the first run there is a total population of 538 agents (nz means non-zero elements) randomly distributed on the sugarscape. As one would expect, during each run some agents die and others move toward the peaks of the sugarscape. For this experiment, the average metabolism of the population goes from 3.5 in the first run to 2 in the sixth run while the average vision goes from 3.5 to 3.8. Thus, as one should expect, lower metabolism and higher vision increase the chances of survival. We can see also that the population tends to reach a stable size and spatial configuration.

Figure 14.7 below shows the carrying capacity of the sugarscape - that is what population size the sugarscape can support - as a function of the maximum level of vision and metabolism of the agents. For each level of vision and metabolism, the average value of ten simulations of six runs each is presented. We can observe how a larger vision and a smaller metabolism tend to increase the carrying capacity of the sugarscape.




Figure 14.7 Carrying Capacity

5. Experiments

A simple experiment would be to add moving cost proportional to the distance moved. This will tend to slow down the convergence to the hilltop locations.

Also, the Sugarscape model can be extended in a number of ways so that many experiments of increasing grade of complexity can be performed. A first step in that direction would be to replace the Sugarscape rule G_{∞} used above with the following one:

Sugarscape growback rule G_1 : At each lattice position, sugar grows back at a rate of α units per time interval up to the capacity at that position.

To introduce this rule, you may want to start by transforming the sugarscape matrix s into a structure with two fields, one containing the capacity and the other the current level of sugar. Then, you can check how different growback rules affect the results.

You may also try to work with agents with finite lives, where their maximum age is a random integer drawn from a given interval $[a,b]$. Then, you may introduce an agent replacement rule such as the following one

Agent replacement rule $R_{[a,b]}$: When an agent dies it is replaced by an agent of age 0 having random genetic attributes, random position on the sugarscape, random initial endowment, and a maximum age randomly selected from the range $[a,b]$.

Epstein and Axtell (1996) present a number of rules for pollution formation, agent mating, agent inheritance, trade, credit, etc., that can be implemented in the Sugarscape model. To learn about the specifics of these rules you are referred to their book.

6. Further Reading

For a comprehensive presentation of the Sugarscape model see Epstein and Axtell (1996). See also the web page of the Sugarscape model at the Brookings Institution at www.brook.edu/es/dynamics/sugarscape/default.htm . For an online guide to agent-based modeling see Axelrod and Tesfatsion (2004). For an approach to estimating agent based models see Gilli and Winker (2003).

For a recent conference keynote address on agent based modeling and an application to finance see LeBaron (2004). Also see his survey paper on agent based computational finance (LeBaron (2005)) which will appear in the Judd and Tesfatsion (2005) volume containing many state-of-the-art papers on agent based modeling. For a comprehensive site with resources on Agent-Based Computational Economics, see the web site developed by Leigh Tesfatsion at www.econ.iastate.edu/tesfatsi/ace.htm . For a review of agent-based modeling as an approach to economic theory, see Tesfatsion (2005).